

EDA Case Study:Houses

By Jacob Alamparambil

Overview

- Preprocessing stage
- mls
- Lot_acres
- Sqr_ft
- Hoa
- Year_built
- Kitchen features
- Garage null values
- Bathroom null values
- Taxes
- statistics

Preprocessing Stage

- Load csv file
- df.head()
- [df.info\(\)](#)
- df.describe()

Remove duplicate records

- 1) Identify primary key : mls

```
df.iloc[0:,0]
```

	MLS
0	21530491
1	21529082
2	3054672
3	21919321
4	21306357
...	...
4995	21810382
4996	21908591
4997	21832452
4998	21900515
4999	4111490

```
df.iloc[0:,0]
```

Remove duplicate records

- 2) Display that column

MLS	
0	21530491
1	21529082
2	3054672
3	21919321
4	21306357
...	...
4995	21810382
4996	21908591
4997	21832452
4998	21900515
4999	4111490

5000 rows × 1 columns

dtype: int64

Remove duplicate records

3) Store the keys in a set and compare with actual column of keys

```
mls=set(df.iloc[0:,0])
```

```
len(mls)
```

```
5000
```



```
len(df.iloc[0:,0])
```



```
5000
```

Change lot_acres null values to median

Before

```
print(df.isnull().sum())
```

MLS	0
sold_price	0
zipcode	0
longitude	0
latitude	0
lot_acres	10
taxes	0
year_built	0
bedrooms	0
bathrooms	6
sqrt_ft	56
garage	7
kitchen_features	33
fireplaces	0
floor_covering	1
HOA	562

lot_acres

0.99

0.99

0.99

0.99

After

```
print(df.isnull().sum())
```

MLS	0
sold_price	0
zipcode	0
longitude	0
latitude	0
lot_acres	0
taxes	0
year_built	0
bedrooms	0
bathrooms	6
sqrt_ft	56
garage	7
kitchen_features	33
fireplaces	0
floor_covering	1
HOA	562

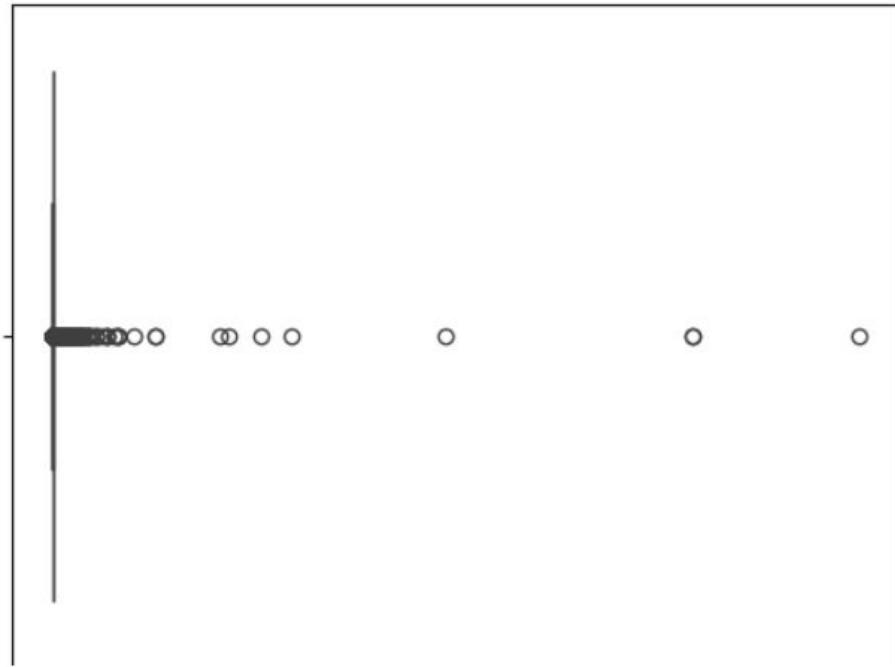
Lot_acres view



```
sns.boxplot(x=df['lot_acres'])
```



```
<Axes: xlabel='lot_acres'>
```



Why?

The data is skewed right to the lower values and contains many outliers so the appropriate imputation would be the median

Change `sqr_ft` null values to median

Why?

The data is skewed right to the lower values and contains many outliers so the appropriate imputation would be the median

Before

```
▶ print(df.isnull().sum())
```

MLS	0
sold_price	0
zipcode	0
longitude	0
latitude	0
lot_acres	0
taxes	0
year_built	0
bedrooms	0
bathrooms	6
sqrt_ft	56
garage	7
kitchen_features	33
fireplaces	0
floor_covering	1
HOA	562

`sqrt_ft`

3512.0
3512.0
3512.0
3512.0

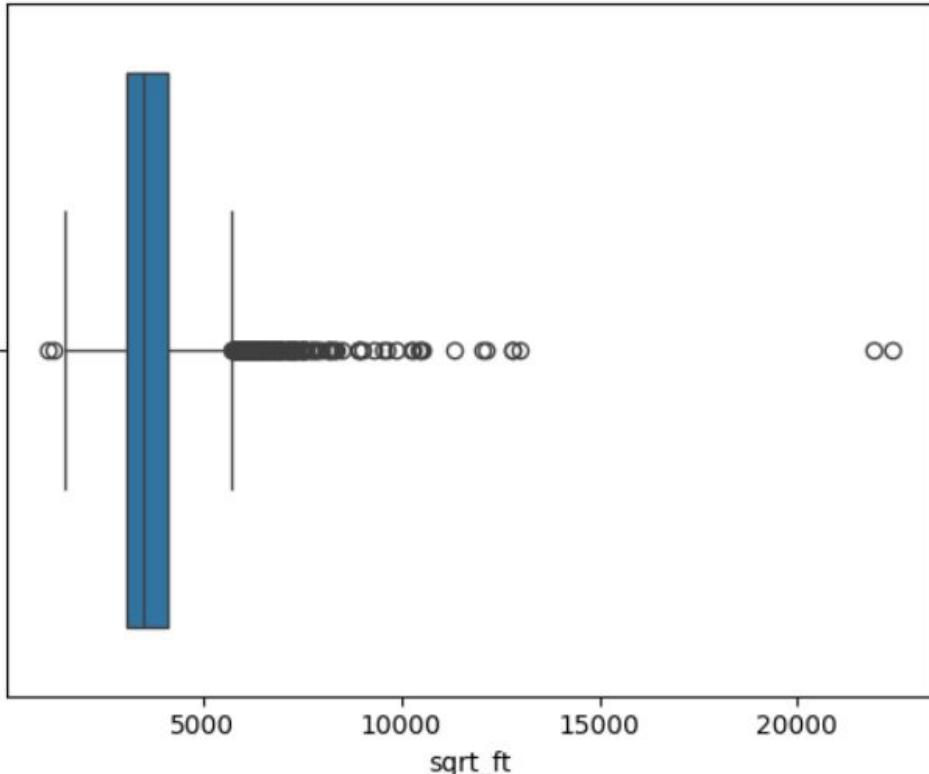
After

```
▶ print(df.isnull().sum())
```

MLS	0
sold_price	0
zipcode	0
longitude	0
latitude	0
lot_acres	0
taxes	0
year_built	0
bedrooms	0
bathrooms	6
sqrt_ft	0
garage	7
kitchen_features	33
fireplaces	0
floor_covering	1
HOA	562

sqr_ft view

<Axes: xlabel='sqrt_ft'>



Why?

The data is skewed right to the lower values and contains many outliers so the appropriate imputation would be the median

Change hoa null value:None to median

Why is hoa=0 ok?

```
▶ print(df.isnull().sum())
```

```
MLS          0
sold_price   0
zipcode      0
longitude    0
latitude     0
lot_acres    0
taxes        0
year_built   0
bedrooms     0
bathrooms    6
sqrt_ft      0
garage       7
kitchen_features 33
fireplaces   0
floor_covering 1
HOA          562
dtype: int64
```

```
▶ print(df.isnull().sum())
```

```
MLS          0
sold_price   0
zipcode      0
longitude    0
latitude     0
lot_acres    0
taxes        0
year_built   0
bedrooms     0
bathrooms    6
sqrt_ft      0
garage       7
kitchen_features 33
fireplaces   0
floor_covering 1
HOA          0
dtype: int64
```

Change properties with \$0 taxes to median value

```
df.iloc[100:,:]
```

	MLS	sold_price	zipcode	longitude	latitude	lot_acres	taxes	year_built	bedrooms	b
100	21522589	2100000.0	85718	-110.885186	32.327907	1.55	15668.99	2009	5	
101	21814436	1450000.0	85624	-110.738083	31.533643	164.30	0.00	1997	1	
102	21326674	1700000.0	85750	-110.846362	32.329630	1.12	17898.00	2002	4	

```
df.iloc[100:,:]
```

	MLS	sold_price	zipcode	longitude	latitude	lot_acres	taxes	year_built	bedrooms	b
100	21522589	2100000.0	85718	-110.885186	32.327907	1.55	15668.99	2009	5	
101	21814436	1450000.0	85624	-110.738083	31.533643	164.30	6223.76	1997	1	
102	21326674	1700000.0	85750	-110.846362	32.329630	1.12	17898.00	2002	4	

Change year built values to median

```
df['year_built']=years
```



```
years= [years_rep if x == 0 else x for x in list(df['year_built'])]
```

```
len(df[df['year_built']==years_rep])
```

169

```
df['year_built']!=0
```

	year_built
0	False
1	False
2	False
3	False
4	False
...	...
4995	False
4996	False
4997	False
4998	False
4999	False

5000 rows × 1 columns

Change kitchen features null values to mode

```
print(df.isnull().sum())
```

MLS	0
sold_price	0
zipcode	0
longitude	0
latitude	0
lot_acres	0
taxes	0
year_built	0
bedrooms	0
bathrooms	0
sqrt_ft	0
garage	0
kitchen_features	0
fireplaces	0
floor_covering	1
HOA	0
dtype: int64	

Change garage null values to median

```
▶ print(df.isnull().sum())
```

```
MLS          0  
sold_price   0  
zipcode      0  
longitude    0  
latitude     0  
lot_acres    0  
taxes        0  
year_built   0  
bedrooms     0  
bathrooms    6  
sqrt_ft      0  
garage       7  
kitchen_features 33  
fireplaces   0  
floor_covering 1  
HOA          0  
dtype: int64
```

```
▶ print(df.isnull().sum())
```

```
MLS          0  
sold_price   0  
zipcode      0  
longitude    0  
latitude     0  
lot_acres    0  
taxes        0  
year_built   0  
bedrooms     0  
bathrooms    6  
sqrt_ft      0  
garage       0  
kitchen_features 33  
fireplaces   0  
floor_covering 1  
HOA          0  
dtype: int64
```

Roundup garage values to nearest whole number

`garage` |

3.0

`garage`

3

3.0

3

4.5

5

3.0

3

Change bathroom null values to median

```
▶ print(df.isnull().sum())
```

```
MLS                 0
sold_price          0
zipcode             0
longitude           0
latitude            0
lot_acres           0
taxes               0
year_built          0
bedrooms            0
bathrooms           6
sqrt_ft              0
garage               0
kitchen_features    33
fireplaces           0
floor_covering      1
HOA                  0
dtype: int64
```

```
▶ print(df.isnull().sum())
```

```
MLS                 0
sold_price          0
zipcode             0
longitude           0
latitude            0
lot_acres           0
taxes               0
year_built          0
bedrooms            0
bathrooms           0
sqrt_ft              0
garage               7
kitchen_features    33
fireplaces           0
floor_covering      1
HOA                  0
dtype: int64
```

Change properties with \$0 taxes to median value pt2

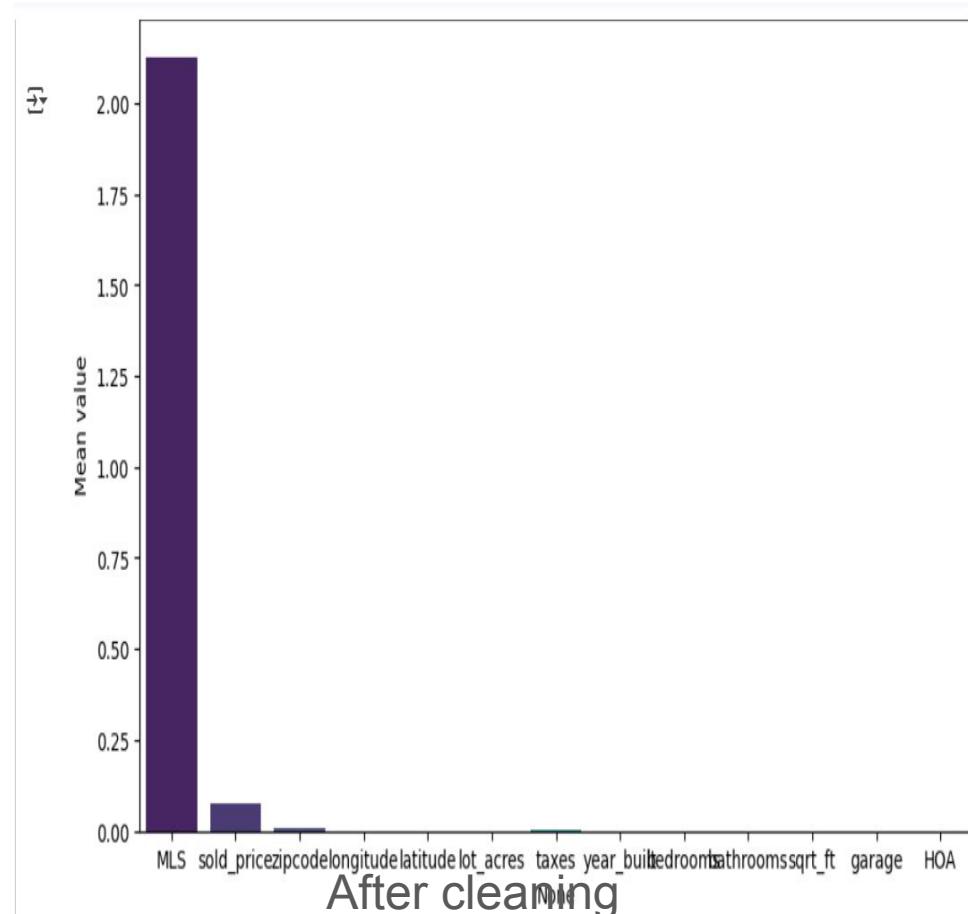
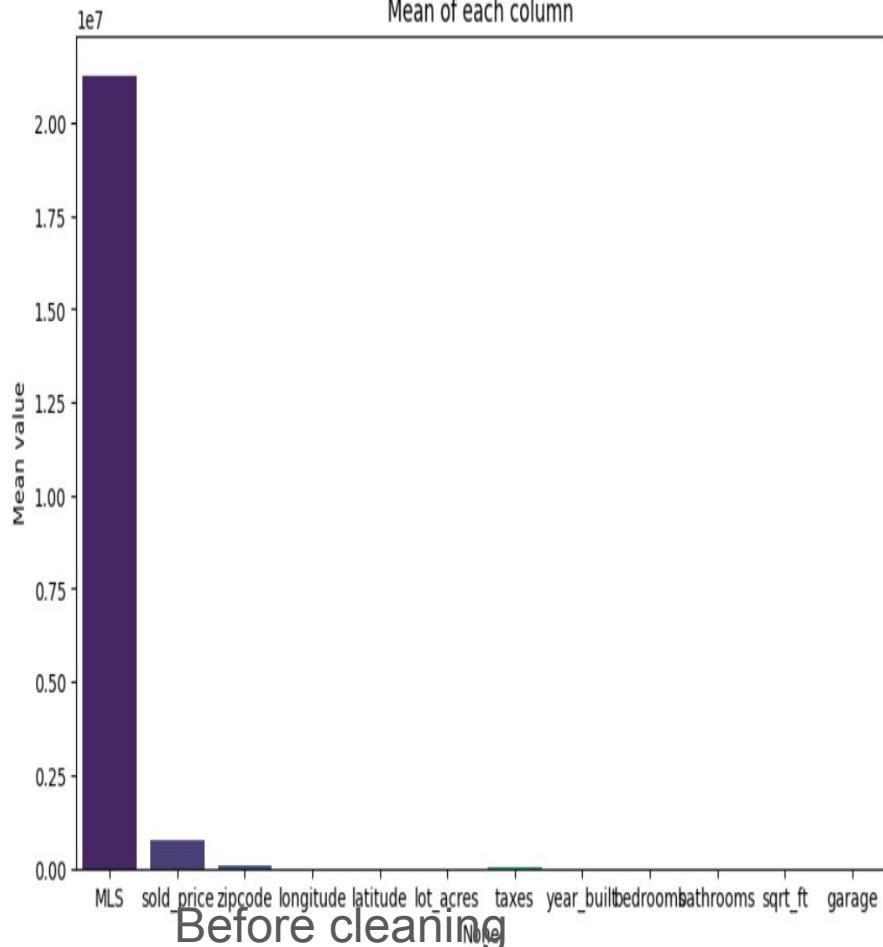
Why?

The data is skewed right to the lower values and contains many outliers so the appropriate imputation would be the median

Comparisons of clean vs unclean data(mean)

index	mean(dirty)	mean(clean)	% Difference
MLS	21270699.22	21270699.22	0
sold_price	774626.2023	774626.2023	0
zipcode	85723.0256	85723.0256	0
longitude	-110.9121066	-110.9121066	0
latitude	32.30851202	32.30851202	0
lot_acres	4.661316633	4.653974	0.1575227254
taxes	9402.828094	9430.212638	0.2912373142
year_built	1992.328	1992.328	0
bedrooms	3.9338	3.9338	0
bathrooms	3.829895875	3.8301	0.005329779364
sqrt_ft	3716.366828	3714.07792	0.06158993732
garage	2.8161426	2.8232	0.2506052073
	Average difference:		0.0638570803

Comparisons of clean vs unclean data(mean pt2)

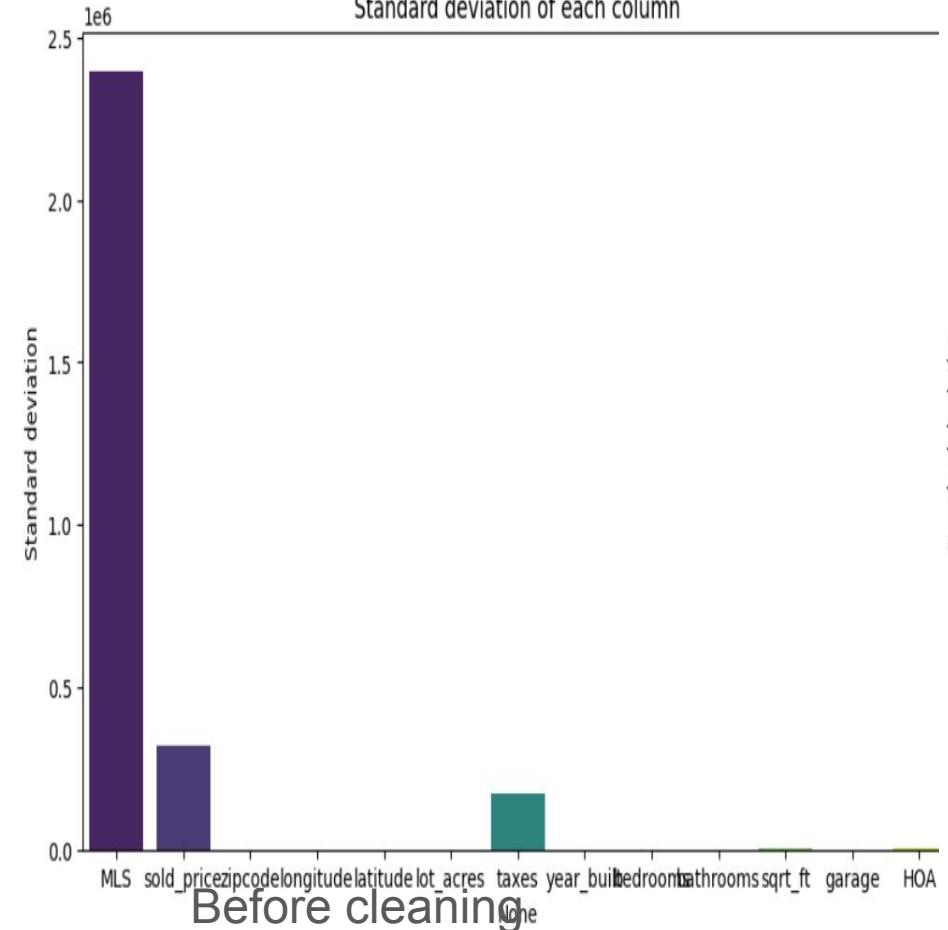


Comparisons of clean vs uncleaned data(std)

index	std(dirty)	std(clean)	% Difference
MLS	2.40E+06	2.40E+06	0.00E+00
sold_price	3.19E+05	3.19E+05	0.00E+00
zipcode	38.061712	38.061712	0
longitude	0.120629	0.120629	0
latitude	0.178028	0.178028	0
lot_acres	51.68523	51.633769	0.09956616233
taxes	1.73E+05	1.73E+05	5.78E-04
year_built	65.48614	65.48614	0
bedrooms	1.245362	1.245362	0
bathrooms	1.387063	1.386243	0.05911771852
sqrt_ft	1120.683515	1114.596305	0.5431694068
garage	1.192946	1.193745	0.06697704674
	average difference:		6.41E-02

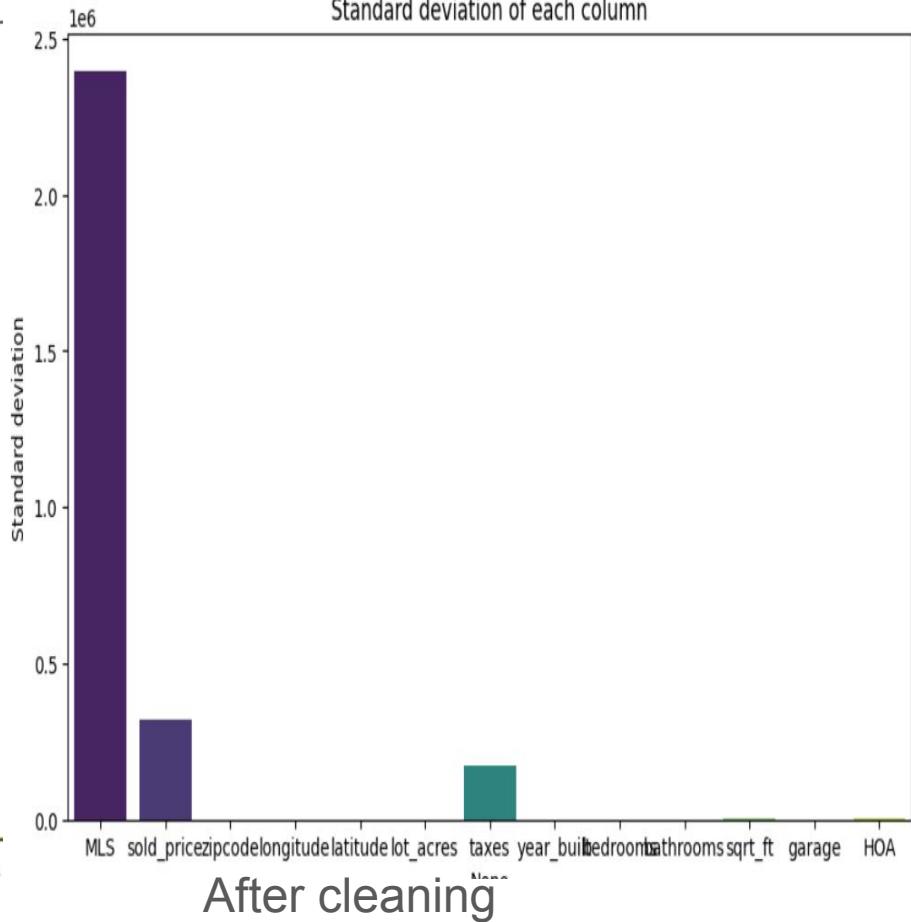
Comparisons of clean vs uncleaned data(std pt2)

Standard deviation of each column



Before cleaning

Standard deviation of each column



After cleaning

Any Questions???